

Introduction to the Apply Function Family

Revolution Analytics





- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`



Overview

In this session we use `apply` functions on elements of lists.

Objectives:

- Understand what are the `apply()` functions
- Understand how `apply()` and family simplifies your code
- Understand different uses of the `apply()` functions





Outline

- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`





Example Dataset

```
head(mtcars)
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1





Summary statistics for multiple variable

- Imagine you want to compute a mean for each column in a data frame.
- Traditional programming languages: Write a loop. For example:

```
cars.means <- numeric(ncol(mtcars))  
for (i in 1:ncol(mtcars)) {  
  cars.means[i] <- mean(mtcars[, i])  
}  
cars.means
```

```
## [1] 20.0906 6.1875 230.7219 146.6875 3.5966 3.2172 17.8487  
## [8] 0.4375 0.4062 3.6875 2.8125
```





A simpler way by using `apply()`

The function `apply()` allows you to compute on the dimensions (e.g. rows and columns) an array or dataframe

```
apply(X, MARGIN, FUN, ...)
```

X an array, including a matrix.

MARGIN a vector giving the subscripts which the function will be applied over. For a matrix, 1 indicates rows, 2 indicates columns.

FUN the function to be applied: see 'Details'. In the case of functions like `+`, `%*%`, etc., the function name must be backquoted or quoted.

optional arguments to FUN.



Looping Across Columns

Remember that the dimensions of a table (and data frame) are:

- 1 Rows
- 2 Columns

So we apply the function `mean()` (or others) to the second dimension of `mtcars`:

```
apply(mtcars, 2, mean)
```

```
##      mpg      cyl      disp      hp      drat      wt      qsec      vs
## 20.0906  6.1875 230.7219 146.6875  3.5966  3.2172 17.8487  0.4375
##      am      gear      carb
## 0.4062  3.6875  2.8125
```




Exercise

Your turn:

- Use `apply()` to calculate the median of each column in the `mtcars` dataset
- Use `apply()` to calculate the 75th %ile of each column (Hint: `?quantile`)





apply() vs. loops

See `help(apply)` for function documentation.

Note the case of `c(1,2)` for MARGIN

```
set.seed(42)
dims <- c(360, 720, 120)
mat.stack <- array(rnorm(prod(dims)), dim = dims)
avgs.mat1 <- matrix(NA, nrow(mat.stack), ncol(mat.stack))
avgs.mat2 <- avgs.mat1
```



apply() vs. loops (cont'd)

```
system.time({  
  for (i in 1:nrow(mat.stack)) {  
    for (j in 1:ncol(mat.stack)) {  
      avgs.mat1[i, j] <- mean(mat.stack[i, j, ])  
    }  
  }  
})
```

```
##      user  system elapsed  
##  6.636   0.008   6.642
```

```
system.time({  
  avgs.mat2 <- apply(mat.stack, c(1, 2), mean)  
})
```

```
##      user  system elapsed  
##  4.804   0.104   4.910
```



Using apply()

Note that the actual data sets are identical.

All differences are 0.

```
table(avgs.mat1 - avgs.mat2)
```

```
##  
##      0  
## 259200
```





Exercise: Calculate `min()` and `max()`

Your turn:

- Use `apply()` to calculate the minimum and maximum values of each column in `mtcars`.
- Use `apply()` to calculate the 90th percentile for each variable in `mtcars` (see `quantile`)





Outline

- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`





Using `lapply()`

The function `lapply()` applies a function to each element of a list, returning the results in list format.

It can be used very much like `apply()` for columns:

```
mins.list.lapply <- lapply(mtcars, min)
mins.list.apply <- apply(mtcars, 2, min)
```





lapply() Results

```
mins.list.lapply
```

```
## $mpg
## [1] 10.4
##
## $cyl
## [1] 4
##
## $disp
...
```

```
mins.list.apply
```

```
##      mpg      cyl  disp    hp  drat    wt   qsec    vs    am  gear
## 10.400  4.000 71.100 52.000  2.760  1.513 14.500  0.000  0.000  3.000
##      carb
##      1.000
```





Why do both work?

```
is.list(mtcars)
```

```
## [1] TRUE
```

```
is.matrix(mtcars)
```

```
## [1] FALSE
```

Both give the same output, but in a different structure





Different data structures

- `apply()` only works on matrices, so it coerces the dataset to a matrix first.
- `lapply()` works on lists

OK sometimes, but you must watch out!





Outline

- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`





Add a character Variable

```
head(mtcars2)
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1   4    4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1   4    4
## Datsun 710      22.8   4  108  93 3.85 2.320 18.61  1  1   4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44  1  0   3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0   3    2
## Valiant        18.1   6  225 105 2.76 3.460 20.22  1  0   3    1
##           efficient
## Mazda RX4      efficient
## Mazda RX4 Wag  efficient
## Datsun 710      efficient
## Hornet 4 Drive  efficient
## Hornet Sportabout gas-guzzler
## Valiant        gas-guzzler
```



Unexpected results with `apply()`

```
apply(mtcars2, 2, mean)
```

##	mpg	cyl	disp	hp	drat	wt	qsec
##	NA	NA	NA	NA	NA	NA	NA
##	vs	am	gear	carb	efficient		
##	NA	NA	NA	NA	NA		



lapply() still works!

```
lapply(mtcars2, mean)
```

```
## $mpg  
## [1] 20.09  
##  
## $cyl  
## [1] 6.188  
##  
## $disp  
... 
```





Why does apply fail?

```
apply(mtcars2, 2, is.numeric)
```

##	mpg	cyl	disp	hp	drat	wt	qsec
##	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
##	vs	am	gear	carb	efficient		
##	FALSE	FALSE	FALSE	FALSE	FALSE		



Why does apply fail?

```
apply(mtcars2, 2, is.character)
```

##	mpg	cyl	disp	hp	drat	wt	qsec
##	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
##	vs	am	gear	carb	efficient		
##	TRUE	TRUE	TRUE	TRUE	TRUE		

Remember: matrices can only store 1 mode of information, so it coerces all values to character.





Outline

- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`





lapply() example

The function `lapply()` returns a list of the same length as `X`, each element of which is the result of applying `FUN` to the corresponding element of `X`.

```
test.frame <- data.frame(X = rnorm(20), Y = rnorm(20, 2, 10), Z = rnorm(20,  
  2, 3))  
test.results <- apply(test.frame, 2, t.test) # apply t-test to columns of the dataframe  
  
class(test.results)  
  
## [1] "list"  
  
length(test.results)  
  
## [1] 3
```



How to extract the names of each result?

Traditional approach:

```
test.names1 <- list()
for (i in 1:length(test.results)) {
  test.names1[[i]] <- names(test.results[[i]])
}
names(test.names1) <- names(test.results)
test.names1
```

```
## $X
## [1] "statistic"      "parameter"      "p.value"        "conf.int"       "estimate"
## [6] "null.value"     "alternative"     "method"         "data.name"
##
## $Y
## [1] "statistic"      "parameter"      "p.value"        "conf.int"       "estimate"
## [6] "null.value"     "alternative"     "method"         "data.name"
...

```





lapply() approach

```
test.names <- lapply(test.results, names)
test.names
```

```
## $X
## [1] "statistic"      "parameter"      "p.value"        "conf.int"       "estimate"
## [6] "null.value"     "alternative"     "method"         "data.name"
##
## $Y
## [1] "statistic"      "parameter"      "p.value"        "conf.int"       "estimate"
## [6] "null.value"     "alternative"     "method"         "data.name"
...
```

```
length(test.names)
```

```
## [1] 3
```





Question:

Can we duplicate `test.names` using `apply()` instead of `lapply()`? If so, what would the `MARGIN` argument be?

```
apply(test.results, 2, names)
```

We cannot easily duplicate `test.names` using `apply()` instead of `lapply()` because `apply()` requires a matrix (i.e. rows and columns).





Extract a Different Element

```
test.pvals <- lapply(test.results, getElement, "p.value")
test.pvals
```

```
## $X
## [1] 0.05243
##
## $Y
## [1] 0.1106
##
## $Z
## [1] 0.0007369
```

```
length(test.pvals)
```

```
## [1] 3
```





Additional helper functions:

sapply() Use `lapply()`, then `simplify2array()`

vapply() Similar to `sapply()`, but you must specify the type of result you expect from each call to FUN. Benefit: tremendously faster than `sapply()`





Exercise:

You turn

- Create a new list of the estimate of each of the `t.test` results contained in `test.results`





Exercise: More practice with lapply

```
data.list <- lapply(sample(0:1, 100, replace = TRUE), function(x) if (x ==  
  0) rnorm(10) else letters[sample(1:26, 1)])  
data.list
```

```
## [[1]]  
## [1] "z"  
##  
## [[2]]  
## [1] "d"  
##  
## [[3]]  
...  
...
```



Exercise: More practice with lapply

- Use `lapply()` on `data.list` to return a list indicating which elements are numeric and which are not.
- For the numeric elements, write a function that returns the average of that element. For the non-numeric elements return the fill value NA.
- Use `sapply()` to do the last step again, except with the output simplified as a vector.





Summary and Questions?

- `apply()` for matrices
- `lapply()` and `sapply()` for lists

Notes: Be careful with `apply()` and `data.frames`!





Outline

- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`





Using by()

The function `by()` is an “object-oriented wrapper” that is applied over subsets of dataframes and matrices defined by the indices.

For example, `mtcars` can be split by row into data frames subset by the levels of `cyl`, and function `FUN` is applied to each subset in turn:

Average of all Variables at each level of Cyl

```
by(mtcars, mtcars$cyl, colMeans)
```

```
## mtcars$cyl: 4
##      mpg      cyl      disp      hp      drat      wt      qsec      vs
## 26.6636  4.0000 105.1364  82.6364  4.0709  2.2857  19.1373  0.9091
##      am      gear      carb
##  0.7273  4.0909  1.5455
## -----
## mtcars$cyl: 6
```



Exercise: Practice with `by()`

Your turn:

- Use `by()` to return summary statistics for all columns of `mtcars` data using the summary function for all combinations of `am` and `vs`.





Outline

- 1 The function `apply()`
- 2 The function `lapply()`
- 3 Example – `lapply()` vs. `apply()`
- 4 `lapply()` example
- 5 The function `by()`
- 6 The function `replicate()`





Using replicate()

The function `replicate()` conveniently evaluates an expression multiple times.

For example, say we want to bootstrap a sample of `1 : 100` ten times:

```
replicate(10, sample(1:100, 100, replace = TRUE))
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## [1,]      52   56   84  100   99   33   89   41   19    99
## [2,]      23   97   19   44   61   89   31   77   83    22
## [3,]      51   31   36   64   27   46   79   23    4    11
## [4,]      48    8   41   56   39   35   46   47    9    60
## [5,]      46    3   18   94    4   11   35   64   80    36
## [6,]      92   17   12   12   95   17   75   87   73    67
## ...
```





Background: The Central Limit Theorem

Central limit theorem (CLT):

- Regardless of the underlying distribution, the distribution of the sum and average of a large number of independent, identically distributed (iid) variables will be approximately normal.





CLT: Background (cont'd)

For large sample sizes,

- The expected value of the sampling distribution of the mean is the population mean: $m_{\bar{X}} = \mu$
- The s.d. of the sampling distribution of the mean is the population: $s_{\bar{X}} = \frac{\sigma}{\sqrt{N}}$

[Useful web page about central limit theorem](#)





Exercise: The CLT and `replicate()`

Your turn:

- Using the sample on the previous slide, generate 100 bootstrap samples of 100 observations from an exponential distribution with a rate parameter of 3.
- Calculate the sample mean of each bootstrap sample.
- Calculate the means and standard deviation of the bootstrapped means. Is this what you would expect?
- Display the results graphically with `hist()`

Hint: Exponential distribution: Use `rexp()`





Module review questions

- What, in a nutshell, is an apply function?
- What are the advantages of using apply functions?
- What are the apply functions we covered in this module, and what does each one do?





Questions?



Thank you

Revolution Analytics is the leading commercial provider of software and support for the popular open source R statistics language.

www.revolutionanalytics.com

1.855.GET.REVO

Twitter: @RevolutionR

