

Data Manipulation with dplyr

Revolution Analytics





- 1 The Grammar of Data Manipulation
- 2 Filtering and Reordering Data
- 3 Transformations and Summaries
- 4 Chaining Verbs
- 5 Advanced Topics





Overview

At the end of this session, you will have learned:

- How to manipulate data quickly with `dplyr` using a very intuitive 'grammar'
- How to use `dplyr` to perform common exploratory and manipulation procedures
- Apply your own custom functions to group manipulations `dplyr` with `mutate()`, `summarise()` and `do()`
- Connect to remote databases to work with larger than memory datasets





Outline

- 1 The Grammar of Data Manipulation
- 2 Filtering and Reordering Data
- 3 Transformations and Summaries
- 4 Chaining Verbs
- 5 Advanced Topics





Why use dplyr?

- R comes with a plethora of base functions for data manipulation
- `dplyr` makes data manipulation easier by providing a few functions for the most common tasks and procedures
- `dplyr` achieves remarkable speed-up gains by using a C++ backend





The dplyr grammar

- dplyr was inspired to give data manipulation a simple, cohesive grammar (similar philosophy to ggplot - grammar of graphics)





Important verbs

filter select rows based on matching criteria

select select columns by column names

arrange reorder rows by column values

mutate add new variables based on transformations of existing variables

summarise reduce variables to smaller values based by groups
(group could be entire dataset)





Newer Versions

Starting with v0.3 (October 2014), there are a host of newer verbs.

slice select rows by number

transmute transform and drop other variables

distinct returns unique rows

count helper verb for tabulating



Data Setup

```
library(dplyr)
dataPath <- "../data"
bankCSV <- file.path(dataPath, "bank-full.csv")
bankData <- read.csv(bankCSV, sep = ";")
```



Viewing Data

- dplyr includes a wrapper called `tbl_df` makes `df` into a 'local `df`' that improves the printing of dataframes in the console
- if you want to see more of the data you can still coerce to `data.frame`

```
(bankData <- tbl_df(bankData))
```

```
## Source: local data frame [45,211 x 17]
##
##   age      job      marital education default balance housing loan
## 1   58  management married  tertiary      no    2143      yes   no
## 2   44  technician single  secondary      no     29      yes   no
## 3   33 entrepreneur married  secondary      no     2      yes  yes
## 4   47  blue-collar married   unknown      no   1506      yes   no
## ...
```



Outline

- 1 The Grammar of Data Manipulation
- 2 Filtering and Reordering Data
- 3 Transformations and Summaries
- 4 Chaining Verbs
- 5 Advanced Topics





Subsetting Data

- `dplyr` makes subsetting by rows very easy
- The `filter` verb takes conditions for filtering rows based on conditions
- This works in a similar fashion to SAS's data step with a `where` clause:





Subsetting Example 1

```
filter(bankData, default == "yes")
```

```
## Source: local data frame [815 x 17]
```

```
##
```

```
##   age      job   marital education default balance housing loan
## 1  42 entrepreneur divorced  tertiary      yes         2     yes   no
## 2  55      services divorced secondary      yes         1     yes   no
## 3  51      admin.   single  secondary      yes        -2     no    no
## 4  33  technician married  secondary      yes        72     yes   no
## ...
```





Subsetting Example 2

```
filter(bankData, balance < 1000)
```

```
## Source: local data frame [30,569 x 17]
```

```
##
```

```
##   age      job   marital education default balance housing loan
## 1  44 technician single secondary    no      29    yes    no
## 2  33 entrepreneur married secondary    no       2    yes   yes
## 3  33      unknown single  unknown    no       1     no    no
## 4  35 management married  tertiary    no     231    yes    no
## ...
```





Subsetting Example 3

```
filter(bankData, month %in% c("april", "may", "jun"), default == "yes")
```

```
## Source: local data frame [351 x 17]
```

```
##
```

```
##   age      job   marital education default balance housing loan
## 1  42 entrepreneur divorced  tertiary     yes         2     yes   no
## 2  55    services divorced secondary     yes         1     yes   no
## 3  51    admin.   single  secondary     yes        -2     no    no
## 4  33 technician married  secondary     yes        72     yes   no
## ...
```





Exercise

Your turn:

- How many defaults in the dataset?
- How many of the entrepreneurs that defaulted were also divorced?





Numeric Indices

You can also extract particular rows by number using `slice()`.

```
slice(bankData, 5:10)
```

```
## Source: local data frame [6 x 17]
```

```
##
```

##	age	job	marital	education	default	balance	housing	loan	contact
## 1	33	unknown	single	unknown	no	1	no	no	unknown
## 2	35	management	married	tertiary	no	231	yes	no	unknown
## 3	28	management	single	tertiary	no	447	yes	yes	unknown
## 4	42	entrepreneur	divorced	tertiary	yes	2	yes	no	unknown

```
...
```





Select a set of columns

- You can use the `select()` verb to specify which columns of a dataset you want
- This is similar to the `keep` option in SAS's data step.
- Use a colon `:` to select all the columns between two variables (inclusive)
- Use `contains` to take any columns containing a certain word/phrase/character





Select Example 1

```
select(bankData, age, job, default, balance, housing)
```

```
## Source: local data frame [45,211 x 5]
##
##   age      job default balance housing
## 1  58 management    no    2143     yes
## 2  44 technician    no     29     yes
## 3  33 entrepreneur  no      2     yes
## 4  47 blue-collar  no   1506     yes
## ...
```





Select Example 2

```
select(bankData, default:duration, contains("p"))
```

```
## Source: local data frame [45,211 x 12]
```

```
##
```

```
##   default balance housing loan contact day month duration campaign pdays
## 1      no    2143     yes  no unknown   5   may      261         1     -1
## 2      no      29     yes  no unknown   5   may      151         1     -1
## 3      no       2     yes  yes unknown   5   may       76         1     -1
## 4      no    1506     yes  no unknown   5   may       92         1     -1
## ...
```





Select: Other Options

`starts_with(x, ignore.case = FALSE)` name starts with x

`ends_with(x, ignore.case = FALSE)` name ends with x

`matches(x, ignore.case = FALSE)` selects all variables whose name matches the regular expression x

`num_range("V", 1:5, width = 1)` selects all variables (numerically) from V1 to V5.

You can also use a `-` to drop variables.



Exercise

Try to use `select()` to

- Drop the first 5 variables in the bank data set.
- Also drop all variables that start with `p`

Hint: Use the `—` to drop variables.





Renaming Variables

There is also a `rename` verb to easily rename variables. You can use `select` to select and to rename by using named arguments.

```
select(bankData, bought_option = y)
```

```
## Source: local data frame [45,211 x 1]
##
##   bought_option
## 1             no
## 2             no
## 3             no
## 4             no
## ...
```





Rename Example

```
rename(bankData, bought_option = y)
```

```
## Source: local data frame [45,211 x 17]
```

```
##
```

```
##   age      job   marital education default balance housing loan
## 1  58  management married  tertiary      no    2143     yes    no
## 2  44  technician  single  secondary      no     29     yes    no
## 3  33 entrepreneur married  secondary      no     2     yes   yes
## 4  47  blue-collar married   unknown      no   1506     yes    no
## ...
```





Reordering Data

- You can reorder your dataset based on conditions using the `arrange()` verb





Arrange Example 1

```
arrange(bankData, job, default)
```

```
## Source: local data frame [45,211 x 17]
```

```
##
```

```
##   age    job   marital education default balance housing loan contact day
## 1  41 admin. divorced secondary    no     270    yes   no unknown   5
## 2  29 admin.   single secondary    no     390    yes   no unknown   5
## 3  45 admin.   single  unknown    no      13    yes   no unknown   5
## 4  44 admin. married secondary    no    -372    yes   no unknown   5
## ...
```





Arrange Example 2

```
arrange(bankData, balance, default)
```

```
## Source: local data frame [45,211 x 17]
```

```
##
```

```
##   age      job      marital education default balance housing loan
## 1  26 blue-collar single secondary    yes   -8019      no    yes
## 2  49  management married  tertiary    yes  -6847      no    yes
## 3  60  management divorced tertiary    no   -4057     yes    no
## 4  43  management married  tertiary    yes  -3372     yes    no
## ...
```





Arrange Example 3

You can use `desc()` to sort in descending order.

```
arrange(bankData, desc(balance), default)
```

```
## Source: local data frame [45,211 x 17]
```

```
##
```

##	age	job	marital	education	default	balance	housing	loan
## 1	51	management	single	tertiary	no	102127	no	no
## 2	59	management	married	tertiary	no	98417	no	no
## 3	84	retired	married	secondary	no	81204	no	no
## 4	84	retired	married	secondary	no	81204	no	no

```
...
```





Exercise

Use `arrange()` to sort on the basis of `y`, `marital`, `job` (in descending order), and `balance`





Summary

filter() Extract subsets of rows. See also `slice()`

select() Extract subsets of columns. See also `rename()`

arrange() Sort your data



Questions?





Outline

- 1 The Grammar of Data Manipulation
- 2 Filtering and Reordering Data
- 3 Transformations and Summaries
- 4 Chaining Verbs
- 5 Advanced Topics





Transformations

- The `mutate()` verb can be used to make new columns

```
mutate(bankData, DefaultFlag = ifelse(default == "yes", 1, 0))
```

```
## Source: local data frame [45,211 x 18]
```

```
##
```

```
##   age      job   marital education default balance housing loan
## 1  58  management married  tertiary      no    2143     yes   no
## 2  44  technician single  secondary     no      29     yes   no
## 3  33 entrepreneur married  secondary     no       2     yes  yes
## 4  47  blue-collar married   unknown     no    1506     yes   no
```

```
...
```





Transformations 2

```
mutate(bankData, BalanceByDuration = balance/duration)
```

```
## Source: local data frame [45,211 x 18]
```

```
##
```

```
##   age      job      marital education default balance housing loan
## 1  58  management married  tertiary      no    2143      yes   no
## 2  44  technician  single secondary      no     29      yes   no
## 3  33 entrepreneur married secondary      no     2      yes  yes
## 4  47  blue-collar married  unknown      no   1506      yes   no
```

```
...
```





Transmute

`mutate()` retains all columns. If you only want to keep the new transforms, you can use `transmute()`





Transmute Example

```
transmute(bankData, BalanceByDuration = balance/duration)
```

```
## Source: local data frame [45,211 x 1]
```

```
##
```

```
##   BalanceByDuration
```

```
## 1      8.210727969
```

```
## 2      0.192052980
```

```
## 3      0.026315789
```

```
## 4     16.369565217
```

```
...
```





Summarise Data by Groups

- The `group_by` verb creates a grouping by a categorical variable
- Functions can be placed inside `summarise` to create summary functions

```
args(group_by)
```

```
## function (.data, ..., add = FALSE)  
## NULL
```





Example group_by 1

```
summarise(group_by(bankData, default), Num = n())
```

```
## Source: local data frame [2 x 2]
```

```
##
```

```
##   default    Num
```

```
## 1      no 44396
```

```
## 2     yes   815
```





Example group_by 2

```
summarise(group_by(bankData, default), Ave.Balance = mean(balance))
```

```
## Source: local data frame [2 x 2]
```

```
##
```

```
##   default Ave.Balance
```

```
## 1      no  1389.8064
```

```
## 2     yes  -137.6245
```



Example group_by 3

```
summarise(group_by(bankData, default), Ave.Balance = mean(balance),  
  Num = n())
```

```
## Source: local data frame [2 x 3]  
##  
##   default Ave.Balance   Num  
## 1      no   1389.8064 44396  
## 2     yes   -137.6245   815
```





Outline

- 1 The Grammar of Data Manipulation
- 2 Filtering and Reordering Data
- 3 Transformations and Summaries
- 4 Chaining Verbs
- 5 Advanced Topics





Chaining/Piping

- A dplyr installation includes the magrittr package as a dependency
- The magrittr package includes a pipe operator that allows you to pass the current dataset to another function
- This makes interpreting a nested sequence of operations much easier to understand





Standard Code

Code is executed inside-out.

```
arrange(filter(select(bankData, age, job, education, default), default ==  
  "yes"), job, education, age)
```

```
## Source: local data frame [815 x 4]
```

```
##
```

```
##   age    job education default
```

```
## 1  25 admin. secondary    yes
```

```
## 2  26 admin. secondary    yes
```

```
## 3  26 admin. secondary    yes
```

```
## 4  26 admin. secondary    yes
```

```
...
```





Reformatted

```
arrange(  
  filter(  
    select(bankData, age, job, education, default),  
    default == 'yes'),  
  job, education, age)
```

```
## Source: local data frame [815 x 4]
```

```
##
```

```
##   age    job education default  
## 1  25 admin. secondary    yes  
## 2  26 admin. secondary    yes  
## 3  26 admin. secondary    yes  
## 4  26 admin. secondary    yes
```

```
...
```





With Piping

```
bankData %>%  
  select(age, job, education, default) %>%  
  filter(default == 'yes') %>%  
  arrange(job, education, age)
```

```
## Source: local data frame [815 x 4]  
##  
##   age    job education default  
## 1   25 admin. secondary    yes  
## 2   26 admin. secondary    yes  
## 3   26 admin. secondary    yes  
## 4   26 admin. secondary    yes  
... 
```





More General Piping

- Piping is not restricted to `dplyr` manipulation tasks

```
x1 <- rnorm(10)
x2 <- rnorm(10)
sqrt(sum((x1 - x2)^2))
```

```
## [1] 5.17407
```

```
(x1 - x2)^2 %>% sum() %>% sqrt()
```

```
## [1] 5.17407
```





Pipe + group_by()

The pipe operator is very helpful for group by summaries

```
bankData %>% group_by(job) %>% summarise(Number = n(), Average.Balance = mean(balance),  
  Number.Defaulted = sum(default == "yes"), Default.Rate = Number.Defaulted/Number)
```

```
## Source: local data frame [12 x 5]
```

```
##
```

```
##           job Number Average.Balance Number.Defaulted Default.Rate  
## 1      admin.   5171      1135.8389             74  0.014310578  
## 2 blue-collar   9732      1078.8267            201  0.020653514  
## 3 entrepreneur  1487      1521.4701             55  0.036987223  
## 4   housemaid  1240      1392.3952             22  0.017741935
```


```
...
```



Pipe and Plot

As a reminder, piping can also be used for non-dplyr functions.

```
library(ggplot2)
bankData %>% filter(job %in% c("management", "technician", "unemployed")) %>%
  group_by(job, marital) %>% summarise(Counts = n()) %>% ggplot() +
  geom_bar(aes(x = job, y = Counts, fill = marital), stat = "identity",
    position = "dodge")
```





Piping: Unique Values

Piping is also very helpful with identifying unique rows. You can also use `distinct()` to identify unique rows and is typically used with `arrange()`.

```
bankData %>% select(job, marital, education, default, housing, loan,
  contact) %>% arrange(job, marital, education, default, housing,
  loan, contact) %>% distinct()
```

```
## Source: local data frame [1,482 x 7]
##
##      job  marital education default housing loan  contact
## 1 admin. divorced primary      no      no  no  cellular
## 2 admin. divorced primary      no      no  no  unknown
## 3 admin. divorced primary      no      no  yes cellular
## 4 admin. divorced primary      no     yes  no  cellular
...

```





Unique Keys

You can specify variables that you only want unique values for.

```
bankData %>% select(job, marital, education, default, housing, loan,  
  contact) %>% arrange(job, marital, education, default, housing,  
  loan, contact) %>% distinct(job, marital, education)
```

```
## Source: local data frame [144 x 7]
```

```
##
```

```
##      job  marital education default housing loan  contact  
## 1 admin. divorced  primary      no      no  no cellular  
## 2 admin. divorced secondary      no      no  no cellular  
## 3 admin. divorced tertiary      no      no  no cellular  
## 4 admin. divorced  unknown      no      no  no cellular
```

```
...
```





Unique Keys

It will keep the first row with those particular key values.

```
bankData %>% select(job, marital, education, default, housing, loan,
  contact) %>% arrange(job, marital, education, desc(default), desc(housing),
  desc(loan), desc(contact)) %>% distinct(job, marital, education)
```

```
## Source: local data frame [144 x 7]
```

```
##
```

```
##      job  marital education default housing loan  contact
## 1 admin. divorced  primary      no      yes  yes  unknown
## 2 admin. divorced secondary    yes      yes  yes  unknown
## 3 admin. divorced  tertiary    yes      no  yes  cellular
## 4 admin. divorced   unknown    no      yes  no  unknown
```

```
...
```



Exercise

Your turn:

- Use the pipe operator to group by job and housing status
- Calculate the counts of observations, and the average and median balance





Summary

mutate() Create transformations

summarise() Aggregate

group_by() Group your dataset by levels

distinct() Extract unique values (frequently used with `arrange()`)

Chaining with the `%>%` operator can result in more readable code.



Questions?





Outline

- 1 The Grammar of Data Manipulation
- 2 Filtering and Reordering Data
- 3 Transformations and Summaries
- 4 Chaining Verbs
- 5 Advanced Topics





Multiple Columns

You can summarise or mutate multiple columns using the same grouping variable.

- `summarise_each` allows you to apply the same summary function to multiple columns
- `mutate_each` also does a similar manipulation for mutate

```
help(summarise_each)
```

Both use the `funcs()` function.





Summarise_each Example

```
bankData %>% group_by(education) %>% summarise_each(funs(mean), balance,  
  duration)
```

```
## Source: local data frame [4 x 3]  
##  
##   education  balance duration  
## 1   primary 1250.950 255.9330  
## 2 secondary 1154.881 258.6858  
## 3   tertiary 1758.416 258.5185  
## 4   unknown 1526.754 257.3139
```





summarise_each Example 2

You can also use multiple functions.

```
bankData %>% group_by(education) %>% summarise_each(funs(min, mean,  
  max), balance, duration)
```

```
## Source: local data frame [4 x 7]  
##  
##   education balance_min duration_min balance_mean duration_mean  
## 1 primary      -2604           0      1250.950      255.9330  
## 2 secondary   -8019           0      1154.881      258.6858  
## 3 tertiary    -6847           2      1758.416      258.5185  
## 4 unknown     -1445           4      1526.754      257.3139  
## Variables not shown: balance_max (int), duration_max (int)
```





mutate_each Example

You can use the `.` to indicate where the variables go in an arbitrary function.

```
bankData %>% group_by(month) %>% select(balance, duration) %>% mutate_each(funs(half = ./2))
```

```
## Source: local data frame [45,211 x 3]
## Groups: month
##
##   month balance duration
## 1    may  1071.5    130.5
## 2    may   14.5     75.5
## 3    may    1.0     38.0
...

```





Additional Helper Functions

- Helper functions `n()` and `count()` count the number of rows in a group
- Helper function `n_distinct(vector)` counts the number of unique items in that vector

```
bankData %>% group_by(job, default) %>% summarise(education_levels = n_distinct(education))
```

```
## Source: local data frame [24 x 3]
## Groups: job
##
##           job default education_levels
## 1      admin.      no                4
## 2      admin.     yes                3
## 3 blue-collar     no                4
## ...
```



tally

`tally()` is a shortcut for counting

```
bankData %>% group_by(job) %>% tally()
```

```
## Source: local data frame [12 x 2]
##
##           job      n
## 1      admin. 5171
## 2 blue-collar 9732
## 3 entrepreneur 1487
## 4   housemaid 1240
## ...
```





Without tally()

```
bankData %>% group_by(job) %>% summarise(n = n())
```

```
## Source: local data frame [12 x 2]
```

```
##
```

```
##           job      n
```

```
## 1      admin. 5171
```

```
## 2  blue-collar 9732
```

```
## 3  entrepreneur 1487
```

```
## 4    housemaid 1240
```

```
...
```





count

`count()` makes it even easier.

```
bankData %>% count(job)
```

```
## Source: local data frame [12 x 2]
##
##       job      n
## 1   admin. 5171
## 2 blue-collar 9732
## 3 entrepreneur 1487
## 4   housemaid 1240
## ...
```





Ranking Variables

In base R, you can use `rank`.

```
args(rank)
```

```
## function (x, na.last = TRUE, ties.method = c("average", "first",  
##       "random", "max", "min"))  
## NULL
```





dplyr Helper Functions

row_number(x) Ties are ordered by row in which the observation is encountered.

ntile(x, n) Rough ranking into n bins

min_rank(x) Ties are assigned the min rank

dense_rank(x) Same as min, but no gaps in ranks

percent_rank(x) min_rank(x) scaled to $[0, 1]$

cume_dist(x) a cumulative distribution function



Rank Examples

```
bankData %>% slice(1:10) %>%  
  transmute(Job = job,  
            jobRankAvg = rank(job),  
            jobRankRow = row_number(job),  
            jobRankMin = min_rank(job),  
            jobRankDense = dense_rank(job),  
            jobRankPerc = percent_rank(job),  
            jobRankCume = cume_dist(job))
```

```
## Source: local data frame [10 x 7]
```

```
##
```

```
##           Job jobRankAvg jobRankRow jobRankMin jobRankDense jobRankPerc  
## 1 management         5.0          4          4          3  0.3333333  
## 2 technician         8.5          8          8          5  0.7777778  
## 3 entrepreneur       2.5          2          2          2  0.1111111  
## 4 blue-collar        1.0          1          1          1  0.0000000  
... 
```



Applying custom functions

- You can also apply your own custom functions using `do()`

```
set.seed(1)
df <- data.frame(houseID = rep(1:10, each = 10), year = 1995:2004,
  price = ifelse(runif(10 * 10) > 0.5, NA, exp(rnorm(10 * 10))))
head(df)
```

```
##   houseID year   price
## 1      1 1995 1.4890017
## 2      1 1996 0.5422509
## 3      1 1997         NA
## 4      1 1998         NA
## 5      1 1999 4.1913535
## 6      1 2000         NA
```



Dragging Down Values

- In our custom dataframe, we have a missing value for some years
- We would like to pull down the previous year's value

```
library(xts)
df %>% group_by(houseID) %>% do(na.locf(.))
```

```
## Source: local data frame [100 x 3]
## Groups: houseID
##
##   houseID year    price
## 1         1 1995 1.4890017
## 2         1 1996 0.5422509
## 3         1 1997 0.5422509
## ...
```



Grouped Tests

```
bankData %>% filter(marital %in% c("married", "single")) %>% group_by(job) %>%  
  do(tTest = t.test(age ~ marital, data = .)) %>% mutate(tTestPVal = get("p.value",  
    tTest), tTestStat = get("statistic", tTest))
```

```
## Source: local data frame [12 x 4]
```

```
## Groups: <by row>
```

```
##
```

```
##           job           tTest      tTestPVal tTestStat  
## 1      admin. <S3:htest> 7.772104e-164 28.547365  
## 2 blue-collar <S3:htest> 8.743035e-318 41.963819  
## 3 entrepreneur <S3:htest> 2.302137e-35 13.732785  
... 
```



Connecting to a database

- dplyr provides capabilities to connect to remote databases
- This is useful in case your data is too large to fit in memory
- We will show how to connect to a remote `sqlite` database





Database Setup

```
library(DBI)
library(RSQLite)
outPath <- "../output"
if (!file.exists(outPath)) dir.create(outPath)
bankSQL <- file.path(outPath, "bank.sqlite")
db <- dbConnect(SQLite(), dbname = bankSQL)
dbWriteTable(conn = db, name = "bankfull", value = bankCSV, sep = ";",
  header = T, overwrite = T)
```

```
## [1] TRUE
```



dplyr and SQL

We connect with `src_sqlite()`

```
args(src_sqlite)
```

```
## function (path, create = FALSE)  
## NULL
```





Connect with dplyr

```
db.new <- src_sqlite(bankSQL)
mytbl <- tbl(db.new, "bankfull")
mytbl %>% group_by(job) %>% tally()

## Source: sqlite 3.8.6 [../output/bank.sqlite]
## From: <derived table> [?? x 2]
##
##           job      n
## 1      "admin." 5171
## 2  "blue-collar" 9732
## 3  "entrepreneur" 1487
## ...
```





Set Operations

`intersect, union, setdiff, setequal`

all have methods for data frames, data tables and SQL db tables

```
args(dplyr::intersect)
```

```
## function (x, y, ...)  
## NULL
```



Joins

There are methods for doing joins (merges) in dplyr

Base tool: `merge()`

dplyr tool: `XXXX_join()`

`help(join)`





Summary

dplyr is a powerful and **fast** package designed to facilitate many common data manipulation problems.

- Subsetting rows (`filter()` and `slice()`) and columns (`select()` and `rename()`)
- Generating transformations (`mutate()` and `transmute()`)
- Aggregation (`summarise()` and `group_by()`)
- Sorting (`arrange()`), unique value identification (`distinct()`), and ranking (`xxx_rank()`)
- Providing SQL connectivity





Questions?



Thank you

Revolution Analytics is the leading commercial provider of software and support for the popular open source R statistics language.

www.revolutionanalytics.com

1.855.GET.REVO

Twitter: @RevolutionR

