

# Applications - Finance

**Revolution Analytics**





1 Black-Scholes-Merton Model

2 Serial Function

3 Using apply

4 Using rExec





# Parallel Finance in R

- Monte Carlo methods used to price stocks
- RNG used to explore pricing space
- Embarrassingly parallel – easily parallelized
- Will look at pricing European options
- This is a parallelization example, not a tutorial on option pricing
- This session borrows heavily from
  - Das, S. and Granger, B. “Financial Applications with Parallel R.”  
Journal of Investment Management, 7(4), 66-77





# Outline

1 Black-Scholes-Merton Model

2 Serial Function

3 Using apply

4 Using rExec





# European Option Pricing

- Pricing based on Black-Scholes-Merton (BSM) models:

$$S(t+h) = S(t)e^{(r-\frac{\sigma^2}{2})h+\sigma e(t)\sqrt{h}}$$

- $e(t) \sim N(0, 1), i.i.d..$
- $S(t)$ : stock price at time  $t$
- $h$ : small delta  $t$  (fraction of year)
- $r$ : risk free interest rate
- $\sigma$ : stock volatility





# More Details

- Will generate multiple paths of stock prices ( $S_j(t)$ ,  $t = 1 \dots T$  for the  $j$ -th path)
- Each path has  $n$  steps making  $h = T/n$  (time step)
- Final price of stock:

$$CallPrice = e^{-rT} \frac{\sum_m \max[0, S_j(T) - K]}{m}$$

- where  $m$  is the number of paths.
- With this data we are ready to build BSM code





# Outline

1 Black-Scholes-Merton Model

2 Serial Function

3 Using apply

4 Using rxExec





# Naive, Serial Code

```
mcoption <- function(s0, k, t, v, r, n, m) {  
  h <- t/n  
  s <- matrix(0, m, n + 1)  
  payoff <- matrix(0, m, 1)  
  for (j in 1:m) {  
    s[j, 1] <- s0  
    for (i in 2:(n + 1)) {  
      s[j, i] <- s[j, i - 1] * exp((r - 0.5 * v^2) * h + rnorm(1) *  
        v * sqrt(h))  
    }  
    payoff[j] <- max(0, s[j, n + 1] - k)  
  }  
  callprice <- mean(payoff) * exp(-r * t)  
}
```





# Your Turn - Brainstorm in group

- What are your observations of the code
- How can you parallelize the BSM function?





# Revising `mcoption()`

- The original code is not very R-like
  - nested for loops recall Fortran
  - payoff variable doesn't need to be a matrix
  - last line uses an assignment that is never used
- A slightly more R-like version uses `apply()` to perform the central computation on each row of the `s` matrix, which we initialize with `s0` rather than `0`.
  - We have to take the transpose of the object returned by `apply` to get back our originally-shaped matrix.
  - We then use `ifelse()` to define payoff as a vector.





# Outline

1 Black-Scholes-Merton Model

2 Serial Function

3 Using apply

4 Using rxExec





# Revising mcoption()

```
mcoption2 <- function(s0, k, t, v, r, n, m) {  
  h <- t/n  
  s <- matrix(s0, m, n + 1)  
  s <- t(apply(s, 1, function(x) {  
    for (i in 2:(n + 1)) {  
      x[i] <- x[i - 1] * exp((r - 0.5 * v^2) * h + rnorm(1) *  
        v * sqrt(h))  
    }  
    x  
  })))  
  
  final <- s[, n + 1]  
  payoff <- ifelse(0 > final - k, 0, final - k)  
  mean(payoff) * exp(-r * t)  
}
```



# Comparing mcoption and mcoption2

```
set.seed(42)
t1 <- system.time(a <- mcoption(100, 100, 1, 0.3, 0.03, 52, 10000))
set.seed(42)
t2 <- system.time(b <- mcoption2(100, 100, 1, 0.3, 0.03, 52, 10000))
all.equal(a, b)
```

```
## [1] TRUE
```

- Excellent! Two functions give same result.



# Near Point of Parallelism

- By recasting the problem as we have done, we can easily see where we can parallelize — the computation of the rows of  $s$ .
- Divide  $m$  over the number of processors, distribute the computation of the rows over the processors, and combine the results.
- Extract the function used to compute the rows so it can be easily re-used.





# The computeRow() function

This function encapsulates the parallelizable part of the algorithm:

```
computeRow <- function(s0, n, r, v, h) {  
  x <- rep(s0, n + 1)  
  for (i in 2:(n + 1)) {  
    x[i] <- x[i - 1] * exp((r - 0.5 * v^2) * h + rnorm(1) * v *  
      sqrt(h))  
  }  
  x  
}
```



# Outline

1 Black-Scholes-Merton Model

2 Serial Function

3 Using apply

4 Using rxExec







# A First Cut at mcparrallel()

- We basically rewrite the middle of mcoption2(), to call computeRow() repeatedly via rxExec():

```
mcparrallel <- function(s0, k, t, v, r, n, m, numw) {  
  h <- t/n  
  s <- unlist(rxExec(computeRow, s0, n, r, v, h, timesToRun = m,  
    taskChunkSize = m/numw))  
  s <- matrix(s, nrow = m, ncol = n + 1, byrow = TRUE)  
  final <- s[, n + 1]  
  payoff <- ifelse(0 > final - k, 0, final - k)  
  mean(payoff) * exp(-r * t)  
}
```



# Remember

- Remember that `rxExec()` always returns a list
- To get a vector of values, you must use `unlist()`.
- To return us to the format `mcoption2()` was using, we call `matrix()` to put the returned vector into a matrix.





# Running mcpParallel()

- mcpParallel() returns the solution
- Need performance data (for this exercise)

```
rxSetComputeContext(RxLocalParallel())
rxOptions(numCoresToUse = 8)
numw <- 8
set.seed(42)
t3 <- system.time({
  mcpParallel(s0 = 100, k = 100, t = 1, v = 0.3, r = 0.03, n = 52,
    m = 10000, numw)
})
```



# Time Three Ways

- `mcoption()` by itself (sequential version)
- `mcoption2()` by itself (sequential version)
- `mcparallel()` with `numw = 4`





# Lab for Session 3

- Can we improve our parallel result?
- Our parallel result in many ways returns us to the nested for loop
- Would be nice to take advantage of the efficiency of `apply()` for each chunk - return a list of matrices, rather than a list of row vectors.





# Second try at parallelizing

```
computeRow2 <- function(s0, n, r, v, h, chunksize) {  
  s <- matrix(s0, nrow = chunksize, ncol = n + 1)  
  s <- t(apply(s, 1, function(x, n, r, v, h) {  
    for (i in 2:(n + 1)) {  
      x[i] <- x[i - 1] * exp((r - 0.5 * v^2) * h + rnorm(1) *  
        v * sqrt(h))  
    }  
    x  
  }, n = n, r = r, v = v, h = h))  
  s  
}
```





# Second try at parallelizing

```
mcparallel2 <- function(s0, k, t, v, r, n, m, numw) {  
  h <- t/n  
  chunksize <- m/(2 * numw)  
  s <- rxExec(computeRow2, s0 = s0, n = n, r = r, v = v, h = h,  
    chunksize = chunksize, timesToRun = m/chunksize)  
  s <- do.call("rbind", s)  
  final <- s[, n + 1]  
  payoff <- ifelse(0 > final - k, 0, final - k)  
  mean(payoff) * exp(-r * t)  
}
```



# Running mcpipeline12()

```
rxSetComputeContext(RxLocalParallel())
rxOptions(numCoresToUse = 8)
numw <- 8
set.seed(42)
t4 <- system.time({
  out2 <- mcpipeline12(s0 = 100, k = 100, t = 1, v = 0.3, r = 0.03,
    n = 52, m = 10000, numw)
})
```







# Compare Run-Times

t1

```
##      user  system elapsed
##    4.675    0.014    4.703
```

t2

```
##      user  system elapsed
##    4.846    0.016    4.878
```

t3

```
##      user  system elapsed
##    0.246    0.221   11.964
```

t4

```
##      user  system elapsed
##    0.102    0.181   11.625
```



# Exercise

Can you make it more efficient?





# Parallel Options Pricing Summary

- Implemented Black-Scholes-Merton option pricing
- While this scales reasonably well, this is not necessarily optimal even in R.
- More optimal `mcoption()` code might scale less well
- Decent scaling even with naïve implementation





# Questions?



# Thank you

Revolution Analytics is the leading commercial provider of software and support for the popular open source R statistics language.

[www.revolutionanalytics.com](http://www.revolutionanalytics.com)

1.855.GET.REVO

Twitter: @RevolutionR

